

SAMPLES

SAMPLES — the workflow vs. a single prompt, same test case, same input

Required by the track: *"a video/document that shows the use of workflow for sample test cases as compared to using a single prompt approach with the same test cases."* This is that document. Every quote below is copied verbatim from a real, saved run — runs/21aug/D1_armA.json (naive single prompt), runs/21aug/D1_armB.json (a token-matched chain-of-thought single prompt, same token budget as the workflow, still one call), and runs/21aug/D1_final_seed2.json (the real workflow, arm C, this project's actual submission). Nothing below is paraphrased into a friendlier shape — if a run said something wrong, it is quoted saying it.

The same input, three times

Case D1 (step21drop/cases/D1/input.md): a resident freelancer in India is paid 5,000 USDC by a foreign client, settling at 03:14 IST on Sunday, 28 June 2026 — a date no Indian bank published an official rate for. All three arms below received the identical facts and the identical corpus of law text. The only thing that changes is how many calls are made and what checks run on the answer before it is shown.

	Arm A	Arm B	Arm C — the workflow
Approach	One prompt, no scaffolding	One prompt, chain-of-thought, same token budget as the workflow	Five model calls + four deterministic checks
Model calls	1	1	5 (small ×2, large ×2, a <i>different</i> model for the adversarial check)

	Arm A	Arm B	Arm C — the workflow
Citation verified against real corpus text?	No	No	Yes — every citation, or the conclusion is dropped
M3 citation validity (eval/score.py)	50.0%	40.0%	100.0%
M4 valuation methods named	2/12	2/12	12/12
Valuation dispute reported	₹250 (0.05%)	₹250 (0.05%)	₹47,868.76 (10.19%)
Own answer attacked before publishing?	No	No	Yes — published either way

1. The valuation figure — how big a gap does each approach actually surface?

Arm A picked a single rate outright:

"The income is taxable as professional services income in INR 469,750 (using the TT BR on 2026-06-29)." — certainty: "settled"

One number. No mention that a second candidate date existed with a different rate, no mention that the crypto market itself had a price the entire weekend. The uncertainty this project's real thesis is about — ₹47,868.76 of it — is invisible in Arm A's output. Not disputed. Not seen.

Arm B, given more reasoning room, gets slightly further:

*"The supplier is liable to pay income tax on the receipt of 5,000 USDC, valued at INR 469,750 to 470,000 based on the TT BR on 2026-06-25 or 2026-06-29." — certainty: "inference", spread reported: **₹250, 0.05%***

Better — it names two candidate bank rates instead of picking one silently. But its own valuation.methods array has exactly two entries, both SBI bank rates. It never asks whether the crypto market's own price on that Sunday is a defensible alternative at all. The **real** spread this receipt carries — twelve defensible methods, ₹4,69,750 to ₹5,17,618 — is **about 191 times larger** than what Arm B reports ($₹47,868.76 \div ₹250 = 191.5$). Not because Arm B made an arithmetic error; because nothing in a single prompt forces it to enumerate every defensible method rather than stop at the first two it thought of.

Arm C (⊗ B, node3_valuation.py — plain arithmetic, no model, valuation.methods.minItems enforced by schema):

12 methods. ₹4,69,750.00 → ₹5,17,618.76. Spread: **₹47,868.76 (10.19%)**.

This is not a smarter model finding a bigger number. It is a deterministic enumeration — 2 official dates × 5 market readings × 2 currency-peg assumptions — that a single prompt has no structural reason to perform completely, and two real single-prompt runs on this exact case, above, didn't.

2. The same wrong citation, twice, with nothing to catch it

Both Arm A and Arm B cite **Rule 206(1)** as the provision that both taxes the receipt *and* prescribes its valuation method:

Arm A: *"mandated_by": "ITR2026-RULE-206(1)"* (valuation_method regime) Arm B: *"provision": "ITR2026-RULE-206(1)"* (valuation_method regime)

This is the exact defect this project's own iteration story documents finding and fixing in its real resolvers (ITERATION-STORY.md item 7, DECISION-D50.md): Rule 206's own opening words scope it to income received *"in foreign currency,"* and a virtual digital asset is defined elsewhere as **not** foreign currency. Rule 206 does not reach this receipt at all. Two independent single-prompt runs made this same scope-reach mistake — and neither has anything downstream that could catch it. No citation matcher checks whether Rule 206 actually governs a VDA. No adversarial pass questions it. It simply ships as "certainty": "settled" (Arm A) or "inference" (Arm B).

Arm C's real workflow record (D1_final_seed2.json), after the same mistake was found and fixed in its own history (DECISION-D50.md, DECISION-D54.md, DECISION-D55.md):

"regime": "valuation_method", "certainty": "lacuna", citation: **Rule 57**, outcome: *"Rule 57 does not provide a specific method for valuing USDC, a virtual digital asset, under s.92... Rule 206 is not applicable as USDC is not defined as foreign currency under s.2(47A)."*

The workflow doesn't just avoid Rule 206's mistake — its own record shows the *reasoning* for why Rule 206 doesn't reach, in a way a reader can check against the corpus text directly. Neither single-prompt arm shows its reasoning as a checkable citation trail at all.

3. Was the answer attacked before being shown?

Arm A and Arm B: no. There is no step in a single prompt that questions the prompt's own output. Whatever the model said in one pass is what a user would see.

Arm C: 🤖 5 (node5_adversarial.py, a **different model family**, decision D41) ran against the frozen record and published 4 attacks — all 4 landed, including one catching a real, unsupported claim ("no deduction obligation... as... the payer is outside India") that the resolver had invented and that nothing else in the pipeline would have caught. Full text of every attack:

runs/21aug/D1_final_seed2_attack.json, rendered on the actual output page in section 05, "What we tried to break."

4. What this comparison is not claiming

This is one case, three real runs, not a claim that single prompts are always this far off — results.md's "Where we lose" section documents cases where the naive baseline actually **beat** the workflow (D1's gap recall, 20-Aug table: Arm A 75% vs. Arm C 25%), and the workflow's own gap detector scored 0% on three of six cases. The comparison here is specific: on the exact question this project exists to answer — how much does a stablecoin receipt's valuation genuinely disagree, and does the law pick one — a single prompt, run twice with two different prompting strategies, both missed nearly the entire real uncertainty and both repeated the exact same wrong-scope citation that took four separate fix cycles to correct in this project's own workflow — with a fifth instance of the same failure pattern found and disclosed, not chased, the same night (DECISION-D55.md). Full metrics, every case, every arm, including where the workflow loses: [results.md](#).